Algorithmic Bias – The Hidden Danger

- **What is Bias?**: When a model produces systematically prejudiced results against certain groups of people.
- **The Source**: Bias usually enters through the **Training Data**. If your historical data contains human prejudice (e.g., historical loan denials for certain districts), the AI will learn and amplify that prejudice.
- **Types of Bias**:
    - **Selection Bias**: The data doesn't represent the whole population.
    - **Label Bias**: The "Success" metric is flawed.
- **The Goal**: Fairness. A data scientist must audit their data to ensure it doesn't discriminate based on gender, race, or geography.

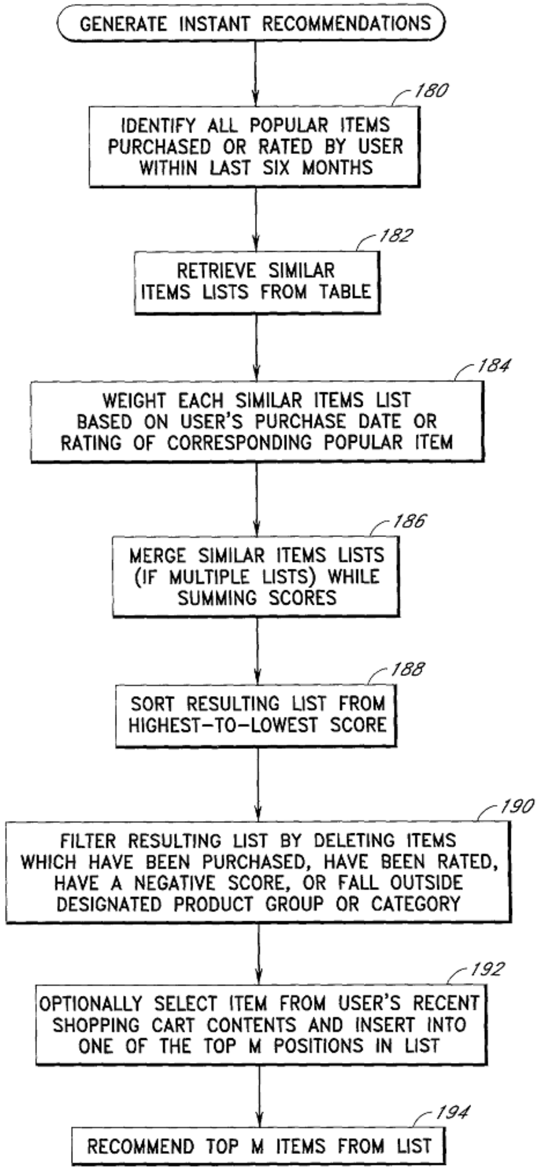

FIG. 5

Data Privacy & Anonymization

- **The Responsibility**: Handling personal information (PII) like names, addresses, or medical IDs requires extreme care.
- **GDPR Standards**: The global benchmark for data protection. It grants users the "Right to be Forgotten."
- **Anonymization Techniques**:
  - **Masking**: Replacing sensitive data with characters (e.g., zaryab@email.com → z****@email.com).
  - **Pseudonymization**: Replacing names with unique IDs.
  - **Aggregation**: Showing "Average Salary per District" instead of individual salaries.

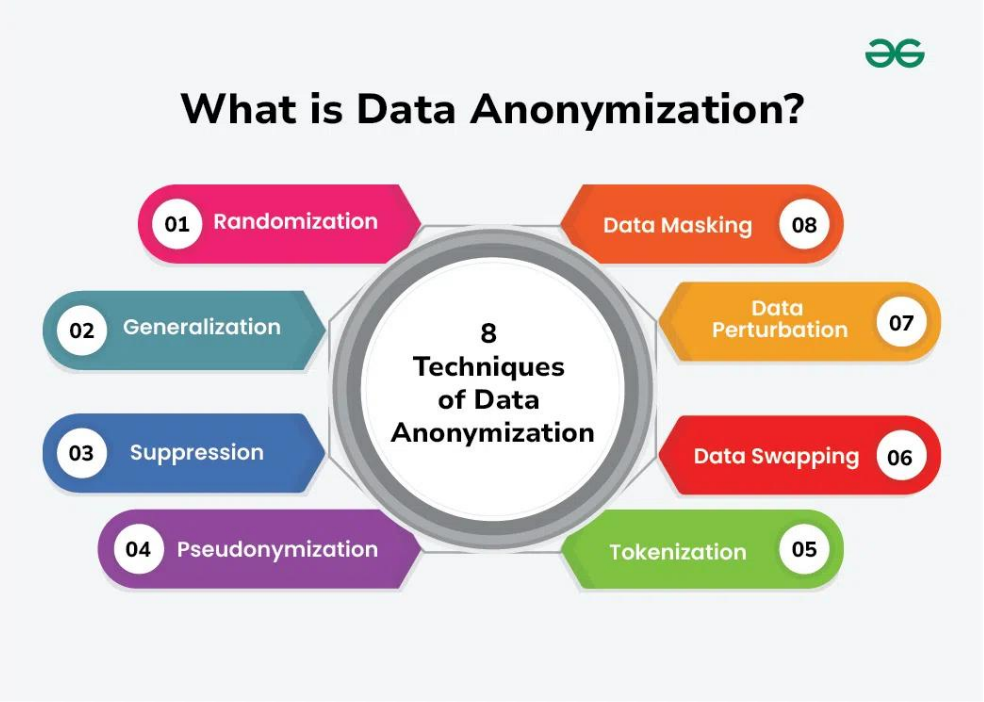

The "Final Pitch" – Speaking to the Board

- **The Language Gap**: CEOs and Managers care about **Value, Risk, and Cost**, not P-values or Hyperparameters.
- **The Structure**:
  - **The Hook**: What problem did we solve?
  - **The Evidence**: Show the most impactful visualization (The "Hero Chart").
  - **The Recommendation**: What should the company do tomorrow?
- **Simplicity**: If you can't explain your model to someone who doesn't know Python, you don't understand your model well enough.