


What is Dimensionality Reduction

Dimensionality Reduction is the process of reducing the number of input variables (features) in a dataset while preserving as much important information as possible.




DIMENSIONALITY REDUCTION

Simplifying Complexity
without Losing the Story


Instructor: Zaryab Ahmad
Khan

THE PROBLEM OF TOO MUCH DATA





- **The "Curse of Dimensionality":** Having 100 columns of data isn't always better.
 - **Performance Hit:** Too many features can confuse an algorithm, increase training time, and actually make predictions *worse* by adding noise.
 - **Analogy:** Trying to recognize a friend by looking at 1,000 microscopic skin cells (too much detail) vs. just looking at their face (the important "components").
-

PRINCIPAL COMPONENT ANALYSIS (PCA)



- **Definition:** A mathematical technique that safely compresses high-dimensional data (e.g., 20 columns) into just a few "Principal Components" (e.g., 3 columns).
- **Goal:** Keep the most important mathematical patterns and "variance" intact while throwing away the redundant fluff.



USING PCA FOR VISUALIZATION

- The Human Limit:** Humans can visualize 2D (a flat chart) or 3D (a cube), but we cannot "see" 50 dimensions.
- The Solution:** Use PCA to shrink complex data down to 2 dimensions so it can be plotted on a standard X/Y scatter chart.
- Benefit:** This allows us to visually confirm if our data naturally forms clusters or groups.